

# Project Report – Neural Network on MNIST Dataset

## 1. Introduction

- Objective: Build and evaluate a neural network to classify handwritten digits (0–9) using the MNIST dataset.
  - Tools: TensorFlow/Keras, scikit-learn, and Keras Tuner for hyperparameter optimization.
  - Dataset: MNIST — 70,000 grayscale images (28×28 pixels).
- 

## 2. Data Preparation

- **Loading Data:** Used `keras.datasets.mnist.load_data()`.
  - **Conversion:** Converted to `float32` and added channel dimension → `(n, 28, 28, 1)`.
  - **Split:**
    - Training: 51,000
    - Validation: 9,000
    - Test: 10,000
  - **Visualization:** Displayed sample digits with their labels.
  - **Preprocessing:**
    - Normalized pixel values.
    - Applied data augmentation for robustness.
- 

## 3. Model Architecture

- Implemented a **Convolutional Neural Network (CNN)** with:
    - Convolution layers (`Conv2D`)
    - Pooling layers (`MaxPooling2D`)
    - Dropout layers for regularization
    - Fully connected dense layers
    - Softmax output with 10 classes
- 

## 4. Hyperparameter Tuning

- Used **Keras Tuner** (RandomSearch) to optimize:
    - Number of filters in Conv2D
    - Kernel sizes
    - Learning rate
    - Dense layer units
  - Best hyperparameters were selected based on validation accuracy.
-

---

## 5. Model Training

- Optimizer: Adam
  - Loss Function: Sparse Categorical Crossentropy
  - Metrics: Accuracy
  - Training strategy:
    - Early stopping
    - Model checkpoints
- 

## 6. Evaluation

- **Classification Metrics:**
    - Accuracy, Precision, Recall, F1-score (via `classification_report`)
  - **Confusion Matrix:** Showed digit-wise misclassifications.
  - **Results:**
    - Training Accuracy: ~99%
    - Validation Accuracy: ~98%
    - Test Accuracy: ~98%
- 

## 7. Insights & Discussion

- The CNN achieved strong generalization with minimal overfitting.
  - Data augmentation improved robustness.
  - Hyperparameter tuning yielded better architecture vs. default setup.
  - Misclassifications were mostly between visually similar digits (e.g., 4 vs 9, 3 vs 5).
- 

## 8. Conclusion

- Successfully built a high-accuracy digit recognition model.
- Validated the effectiveness of CNNs in image classification tasks.
- Future improvements could include:
  - Testing deeper architectures (e.g., ResNet, EfficientNet).
  - Trying advanced optimizers.
  - Deploying the model in a real-time digit recognition application.